

Text

Big data: text

- Text is everywhere. It provides a key source of insight in the modern business environment.
- Examples:
 - Product reviews.
 - Internet searches.
 - Social-media activity surrounding a brand.
 - Company earnings calls.
 - Technical documentation.

Big data: text

- Annoying reality: text data is **unstructured**.
- Compare:

Price	Engine	Cylinders	HP	CityMPG
43755	39014	3.5	6	225
46100	41100	3.5	6	225
36945	33337	3.5	6	265
89765	79978	3.2	6	290
23820	21761	2	4	200
33195	30299	3.2	6	270
26990	24647	2.4	4	200
25940	23508	1.8	4	170

vs.

It is not from the benevolence of the butcher, the brewer, or the baker, that we expect our dinner, but from their regard to their own interest. We address ourselves, not to their humanity but to their self-love, and never talk to them of our own necessities but of their advantages. Nobody but a beggar chooses to depend chiefly upon the benevolence of his fellow-citizens.

Text-analysis goals

- Summary/compression:
 - Representing a body of documents as categories
 - Representing a document with a topic or set of topics
- Find similar documents to a given document
 - Showing users items of similar interest
 - Drawing analogies, e.g. from search terms to ads
- Classification of documents
 - Author attribution

Some terminology

- Corpus/corpora: body of documents
 - Wikipedia (each page is a document)
 - All tweets on a given date (each tweet is a document)
- Dictionary: set of all allowable “words”
- n-gram: set of n words in a row
 - “White House”: a bi-gram
 - See Google’s n-gram browser

Some terminology

- NLP: natural language processing
- Text pipeline: how unstructured text becomes tidy data
- Token versus type
 - Token: a string with an identified meaning.
 - Type: a higher-order category representing a concept.
 - Example: “A rose is a rose is a rose.” — Sylvia Plath
 - 9 tokens, 4 types (“a”, “rose”, “is”, “.”)
- Programming analogy: type = class, token = object.

Some terminology

- tokenization: turning a string of symbols into tokens
- metadata: extra information about a document
 - Examples: author, geo-tag
- XML/JSON: two most common file formats for text data
 - More structured than .txt
 - Includes metadata

Information in a document

- Metadata
- Words in a document
 - Can't describe everything, but can get pretty far
- Sentiment
 - Computers may make mistakes in tagging sentiment
 - Example: I'm having such a great time today at the DMV.
- Main idea/topics
- Grammar/syntax

Tokenization and representation

String of symbols



tokenization

String of tokens



representation

Useful data
structure

All men are dogs



"all", "men", "are", "dogs"
(not: "all m", "en", "ar e",
"dogs")



all: 1
men: 1
are: 1
dogs: 1

Tokenization: from symbols to tokens

- Tokenization involves many choices of what to do and not do to a raw string of symbols.
- Removing/splitting on white spaces
 - Typically the initial step in tokenization
 - Difficult when words are run together (e.e. cummings poem; “onetwothreefourfive”)
- Removing punctuation
 - But be careful, e.g. :-)

Tokenization: from symbols to tokens

- Converting everything to lowercase
- Removing “stop” words
 - the, is, a, of...
 - Take care: one person’s trash is another person’s treasure
- Dropping numbers and mapping to a common symbol (NUM)
- Stemming: drop suffixes
 - acknowledge, acknowledges, acknowledgement
- Deal with misspellings

“Bag of words” (really “bag of types”)

- This shows only the words in a document, and nothing about sentence structure or organization.

“There is a tide in the affairs of men, which taken at the flood, leads on to fortune. Omitted, all the voyage of their life is bound in shallows and in miseries. On such a full sea are we now afloat. And we must take the current when it serves, or lose our ventures.”

- What the data scientist sees:

tide: 1

affairs: 1

men: 1

we: 2

flood: 1

the: 4

fortune: 1

etc.

“Bag of words” (really “bag of types”)

- Advantage: easy to work with and calculate with
- Disadvantage: it destroys other important sources of information (syntax, structure)
- We can get surprisingly far with the BoW representation
- Two common options for data structures:
 - Hash table/key-value store (dictionary in Python)
 - Vectors

“Bag of words” as a hash table

- Keys: words (really types)
- Values: counts in the document
- Easy to add new entry to hash table
- Useful for open-ended vocabulary
- Generally preferred for storage/manipulation

tide: 1	flood: 1
affairs: 1	the: 4
men: 1	fortune: 1
we: 2	...

“Bag of words” as a vector

- Each index in a vector corresponds to a word; each entry is word count
- Can be difficult to update if a new word is encountered
- Easier to do math with
- Two examples of vectors

tide	affairs	the	.	.	.	bacteria	cars	all	men	dogs
1	1	4	.	.	.	0	0	0	1	0
0	0	0	.	.	.	0	0	1	1	1

(most entries not shown)

Document-term matrix (DTM)

- N = number of documents, D = size of dictionary
- X_{ij} = number of times term j appears in document i .
- The document-term matrix X for a corpus puts each document's word vector as a row of the matrix:

(X_{11} X_{12} . . . X_{1D})

(X_{21} X_{22} . . . X_{2D})

(X_{31} X_{32} . . . X_{3D})

. . .

(X_{N1} X_{N2} . . . X_{ND})